

Artist-Conditional Lyric Generation with QLoRA Adapter Blending

Ali Özkaya

Department of Computer Engineering,
Izmir Institute of Technology
aliozkaya00@gmail.com

Abstract. We present a pipeline for artist-conditional lyric generation using per-artist QLoRA adapters on a quantized Gemma 4 E4B model, with inference-time linear blending of adapter weights for style interpolation. We select five thematically similar progressive/extreme metal artists, forcing the model to capture fine-grained stylistic differences rather than genre-level features. A RoBERTa-based artist-attribution classifier (87.3% accuracy, macro-F1 0.86) serves as the primary automated instrument. Adapters reach 0.69–0.99 target-artist attribution with no artist name in the prompt, whereas zero-shot and few-shot baselines collapse onto a single out-of-distribution “sink” artist. Through ablations over adapter method (LoRA vs. DoRA), rank ($r \in \{4, 8, 16\}$), and a style-weighted training loss, together with a cross-artist perplexity matrix and a classifier-free embedding analysis, we show that no single metric captures style: a style-weighted loss can *raise* attribution while degrading text quality, and a 50/50 weight blend is not 50/50 in perceived style. We argue these disagreements are themselves the central methodological finding.

Keywords: lyric generation · QLoRA · adapter blending · parameter-efficient fine-tuning · evaluation methodology

1 Introduction

Large language models excel at open-ended text generation [9], yet controlling stylistic properties remains challenging. In lyric writing, the output must reflect an artist’s distinctive voice—their vocabulary, themes, and rhythmic patterns.

Parameter-efficient fine-tuning methods such as LoRA [6] and QLoRA [3] adapt large models on consumer hardware by training only low-rank weight matrices. A natural extension is training one adapter per artist and combining them at inference time, as explored in adapter merging [12] and adapter soup [1].

We apply this idea to lyric generation, and find that the harder problem is not generation but *evaluation*. Our contributions are:

1. A pipeline for per-artist QLoRA adapters on 4-bit quantized Gemma 4 E4B, using five same-genre artists as a deliberately challenging testbed, with inference-time delta-space adapter blending.

2. A multi-instrument evaluation: a RoBERTa artist-attribution classifier (87.3% accuracy), a cross-artist perplexity matrix, lexical-richness and degeneration metrics, and a classifier-free sentence-embedding analysis.
3. Ablations over LoRA vs. DoRA, rank, and a style-weighted loss, the last of which exposes a single training change that improves one metric while harming others.
4. A demonstration that attribution, lexical breadth, degeneration, and likelihood disagree—*style is not one number*—and an account of *why*.

2 Related Work

Parameter-Efficient Fine-Tuning. LoRA [6] injects a trainable low-rank decomposition $\Delta W = BA$ while freezing pretrained weights. QLoRA [3] adds 4-bit NF4 quantization. DoRA [7] decomposes weight updates into magnitude and direction components.

Creative Generation and Degeneration. LLM-based creative generation spans poetry, storytelling, and songwriting [11]. A recurring failure mode is neural text degeneration—repetitive, looping output produced when low-entropy next-token distributions interact with the model’s copy bias [5]; we observe and quantify this directly as a side effect of style emphasis.

Adapter Composition. LoRAHub [12] learns composition coefficients across task-specific adapters; AdapterSoup [1] averages adapter weights for generalization. We instead linearly interpolate two artist adapters in delta space to study style blending.

Embedding-Based Style Analysis. Sentence encoders such as SentenceBERT [10] map text to a vector space. We repurpose one as a classifier-free attribution instrument by measuring distance to per-artist centroids, giving an axis that was never fit to the artist labels.

3 Methodology

3.1 Pipeline Overview

Our pipeline has four stages: (1) preprocessing, (2) classifier training, (3) per-artist QLoRA adapter training, and (4) inference with optional adapter blending. The classifier validates that real lyrics are distinguishable and later evaluates generated lyrics.

3.2 Preprocessing

Raw lyrics contain structural annotations ([Verse 1], [Chorus]) that encode structure rather than style. We strip all bracketed headers, apply NFKC normalization (folding exotic Unicode whitespace and a Cyrillic homoglyph that had polluted one artist’s distinctive-token vocabulary, while preserving newlines), and filter songs shorter than 100 characters.

3.3 Artist-Attribution Classifier

We fine-tune RoBERTa-base [8] for 5-class classification. The classifier is an evaluation instrument: high-confidence attribution to the target artist indicates successful style capture. We use a maximum sequence length of 512 tokens, covering $\sim 96.7\%$ of songs.

3.4 QLoRA Adapter Training

We use Gemma 4 E4B [4] ($\sim 4\text{B}$ parameters, dense) loaded in 4-bit NF4 quantization with double quantization and bfloat16 compute via bitsandbytes [2]. We chose the dense variant over the MoE one because adapter blending assumes dense layers—merging LoRA weights across routed experts is ambiguous.

LoRA adapters ($r=8$, $\alpha=16$, dropout 0.1) target all attention (q, k, v, o) and MLP (gate, up, down) projections in the language-model layers, selected via regex to exclude the vision/audio encoder layers—the QLoRA-recommended all-linear configuration. The training prompt `"Write song lyrics.\n\n{lyrics}"` is deliberately generic, with *no artist name*, so style is forced into the adapter weights—a prerequisite for clean blending.

3.5 Style-Weighted Loss

As an alternative objective we up-weight each artist’s characteristic vocabulary. Per artist we compute TF-IDF scores over the training corpus, map them to per-token weights w_t (capped at $10\times$, with the mean held near 1 so the overall loss scale is preserved), and minimize a token-weighted cross-entropy:

$$\mathcal{L}_{\text{SW}} = -\frac{1}{\sum_t w_t} \sum_t w_t \log p_\theta(x_t | x_{<t}). \quad (1)$$

Changing the loss changes the learned weights and hence the next-token distribution. Section 6.1 shows that this single change deforms the distribution in two coupled ways, with two distinct symptoms.

3.6 Adapter Blending

At inference time we interpolate two LoRA adapters in delta space by rank concatenation, preserving the low-rank structure exactly:

$$A_{\text{blend}} = \begin{bmatrix} \alpha A_1 \\ (1-\alpha) A_2 \end{bmatrix}, \quad B_{\text{blend}} = [B_1 \ B_2], \quad (2)$$

so that $\Delta W_{\text{blend}} = \alpha \Delta W_1 + (1-\alpha) \Delta W_2$. Endpoints $\alpha \in \{0, 1\}$ recover the pure adapters exactly—a built-in sanity check (numerically verified to $\sim 10^{-7}$).

3.7 Baselines

We compare against the base model with no fine-tuning, under three distinct prompt families (the adapters’ generic prompt is a fourth):

- **B1 – Zero-shot (base):** "Write song lyrics in the style of {artist}."
- **B2 – Few-shot (base):** the same opener, three sampled training lyrics as in-context examples, and a closing cue ("Now write new song lyrics in the style of {artist}:").
- **B3/B4 – Instruct:** the instruction-tuned variant (gemma-4-E4B-it) prompted through its chat template, naming the band *and* genre (to disambiguate names like “Tool”) and requesting lyrics only.

All methods use the same generation parameters (temperature 0.9, top- p 0.9, repetition penalty 1.1, 200–512 tokens) and 10 samples per artist per method; we report mean attribution. Because the four baselines do not share a prompt, cross-baseline differences are partly prompt-driven; we treat the baselines as qualitative motivation rather than a controlled comparison. Crucially, every baseline names the artist while the adapters never do.

4 Experimental Setup

4.1 Dataset and Artist Selection

We draw from the `genius-lyrics-cleaned` dataset on Hugging Face (~ 3.18 M songs) and select five progressive/extreme metal artists who share a genre but differ in voice, so the classifier must learn *artist*, not genre:

- **Mastodon** (154 songs) — sludge/prog, concept-album narratives.
- **Opeth** (118) — poetic, melancholic.
- **Death** (116) — technical, philosophical lyricism.
- **Gojira** (86) — environmental, philosophical.
- **Tool** (74) — cryptic, conversational register.

After preprocessing, 548 songs remain; we apply a stratified 90/10 split (493 training, 55 evaluation).

4.2 Training Configuration

Classifier. RoBERTa-base, 10 epochs, lr 2×10^{-5} , batch 8, weight decay 0.01, warmup ratio 0.1, best checkpoint by validation accuracy (epoch 5).

QLoRA. SFTTrainer (TRL), 3 epochs, effective batch 4, lr 2×10^{-4} , cosine schedule, warmup 0.1, weight decay 0.05, max length 512, gradient checkpointing, on an NVIDIA RTX 5090 (32 GB).

Generation. Nucleus sampling (temperature 0.9, top- p 0.9, repetition penalty 1.1, 200–512 tokens); a fixed seed pairs sampling noise across adapter variants so differences are attributable to the weights.




Fig. 1. Songs per artist.




Fig. 2. Lyric length distributions.

Table 1. Per-class classifier performance on the evaluation set (55 songs). Accuracy 87.3%, macro-F1 0.86, weighted-F1 0.87.

Artist	Precision	Recall	F1
Death	1.00	1.00	1.00
Tool	0.88	1.00	0.93
Mastodon	0.75	1.00	0.86
Opeth	0.90	0.75	0.82
Gojira	1.00	0.56	0.71
Macro avg	0.91	0.86	0.86

4.3 Reproducibility

All code, data splits, cached results, and notebooks are available at <https://github.com/raccoon-42/lora-lyrics> with step-by-step instructions.

5 Results

5.1 Data Analysis

Figure 1 shows moderate class imbalance (Mastodon has $\sim 2\times$ Tool’s songs). Lyric lengths (Figure 2) have a median near 216 tokens; only 3.3% of songs exceed the 512-token limit.

5.2 Classifier Results

Table 1 reports per-class results. Death is perfectly separated; Gojira is weakest (recall 0.56), its misses absorbed mostly by Mastodon—the largest class acts as an in-domain “absorber.” A TF-IDF + logistic-regression baseline reaches 81.8% in 5-fold CV; RoBERTa’s gain confirms it captures deeper patterns than surface lexicon. Validation loss plateaus while training loss keeps falling (Figure 4), indicating mild overfitting. Extending to 10 epochs improved accuracy over a 5-epoch run, but the best checkpoint occurred at epoch 5; since the warmup/decay schedule scales with total steps, the gain is attributable to the *schedule* (a lower effective learning rate at epoch 5 of 10), not to additional epochs.




Fig. 3. Classifier confusion matrix (evaluation set, 55 songs).

5.3 Adapters vs. Baselines, and the Attribution Sink

Figure 5 gives the main result: adapters attribute to the correct artist for all five even though the prompt names none of them—style lives in the weights (representative pure-adapter samples in Appendix A, part A).¹ The baselines fail everywhere except Tool, and that exception is an artifact.

The non-instruction-tuned base model, asked to “write lyrics,” instead continues its pretraining prior for lyrics *web pages* (HTML and forum boilerplate); the classifier has no reject option, so this out-of-distribution text is forced onto the nearest real artist. Every base-model baseline collapses onto **Tool** ($P(\text{Tool})=0.993$ regardless of target), because Tool’s distinctiveness among these five is plain conversational English, which generic web text resembles.

The instruction-tuned model instead produces fluent but mode-collapsed lyrics (near-identical, ornate openings across artists) that collapse onto **Opeth**—whose poetic, melancholic register is the nearest match to a generic “polished poem.” The sink thus follows the *register* of the out-of-distribution input: plain web text routes to Tool, polished verse to Opeth. Two different sinks, *neither* the majority class (Mastodon): the cause is the missing reject option, not class imbalance.

We therefore treat attribution as a sanity check valid only on genuine lyric-like input; adapter output, being lyric-like, is validly scored.




Fig. 4. Classifier training over 10 epochs. Train loss falls to near zero while validation loss plateaus and rises (mild overfitting); validation accuracy peaks at epoch 5 (0.873, the adopted checkpoint).




Fig. 5. Target-artist attribution (mean $\pm 1\sigma$, $n=10$). *Left:* baselines—the base model collapses onto Tool, the instruct model onto Opeth (each scores high only for its sink target). *Right:* adapter ablations, with the y -axis zoomed to expose the LoRA/DoRA/SW differences.

5.4 Ablations: Method and Rank

LoRA vs. DoRA (Figure 5, right panel) is mixed: DoRA wins Gojira and Opeth, LoRA wins Tool, Death, and Mastodon, with no perplexity difference (Section 5.6). Neither dominates at this scale.

Rank (Figure 6, three artists) contradicts the intuitive “ $r=8$ is best”: $r=16$ consistently overfits—lower attribution *and* inflated off-target perplexity—while $r=4$ matches or beats $r=8$ on attribution. Yet on perplexity $r=8$ can win (Tool own-column 4.3 at $r=8$ vs. 8.0 at $r=4$): the best rank depends on the objective, an early instance of the metric disagreement that recurs throughout.

¹ The full set of pure-adapter generations—210 in total (5 artists $\times$ adapter variants $\times$ 10 samples)—is cached as JSON in the repository under `src/results/adapters/`, alongside the per-sample attribution scores.




Fig. 6. LoRA rank ablation (target-artist attribution, mean $\pm 1\sigma$, $n=10$; y -axis zoomed). $r=4$ matches or beats $r=8$ on all three artists, while $r=16$ overfits.

Table 2. Plain vs. style-weighted (SW) adapters, per artist. Attribution = target P ; types and occurrences = markers per sample; distinct-2 lower means more looping. Bold marks the largest moves (Gojira breadth, Tool repetition).

Artist	Attribution		Types		Occurrences		Distinct-2	
	plain	SW	plain	SW	plain	SW	plain	SW
Gojira	.99	.67	1.6	5.7	1.9	11.9	.96	.71
Tool	.69	.90	2.5	3.3	3.1	28.5	.88	.74
Death	.90	.92	1.7	1.5	2.9	7.2	.67	.63
Mastodon	.87	.87	2.0	3.3	2.7	4.9	.90	.70
Opeth	.90	.85	0.7	2.8	0.7	7.1	.92	.71

5.5 Style-Weighted Loss: Metrics Disagree

We score the SW adapters on four axes. For lexical analysis we define a *marker* as one of an artist’s top-40 distinctive tokens (TF-IDF, stopwords dropped); *types*/sample counts distinct markers (breadth), *occurrences*/sample counts total markers (repetition), and *distinct-2* is the fraction of unique bigrams (lower = more looping). Table 2 contrasts plain and SW adapters.

The axes disagree sharply. Gojira’s SW adapter gains genuine breadth ($3.6\times$ more distinct markers) yet *loses* attribution ($.99 \rightarrow .67$): the classifier penalizes the enrichment.

Tool is the opposite—attribution *rises* ($.69 \rightarrow .90$) while breadth stays flat ($2.5 \rightarrow 3.3$) and occurrences explode ($3.1 \rightarrow 28.5$, ~ 8.6 per type). Tool’s gain is therefore pure repetition, not richer vocabulary, and is caught only because the repetition-robust *types* metric did not move—a degenerate Tool “hey” loop scores 0.996 while a coherent verse from the plain adapter scores 0.003 (Appendix A, part C).

Table 3. Cross-artist perplexity (lower better). Each block’s column minimum is bold (specialization). Base row for reference.

adapter ($r=8$)	Gojira	Tool	Death	Mastodon	Opeth
base	45.9	18.6	30.4	32.3	57.7
<i>plain LoRA</i>					
Gojira	18.2	10.9	15.3	18.7	34.0
Tool	28.2	4.3	16.1	19.1	37.0
Death	70.3	24.6	4.5	49.5	105.6
Mastodon	40.4	14.6	23.0	8.1	54.4
Opeth	29.0	12.2	16.5	20.7	9.5
<i>style-weighted</i>					
Gojira	22.2	11.9	17.2	20.2	39.5
Tool	28.1	5.2	15.8	18.7	36.6
Death	49.8	19.1	6.6	36.8	71.3
Mastodon	35.5	13.5	20.7	11.6	47.1
Opeth	35.1	13.4	19.3	22.7	16.9

Distinct-2 falls for all five (degeneration is universal under SW), so it cannot tell the two cases apart. No single metric reveals that Gojira’s SW bought real breadth while Tool’s bought none—and that attribution rewarded the wrong one.

5.6 Perplexity

Perplexity is measured by teacher-forcing each adapter on every artist’s held-out lyrics (the generic prompt masked): the model scores real text it never produced, so it is hard to game by repetition. Table 3 shows own-artist perplexity (column minima, in bold).

Every adapter beats the base model in every column and minimizes its own column, so the adapters specialize on the likelihood axis too. But SW *raises* every diagonal entry (e.g. Gojira $18.2 \rightarrow 22.2$, Opeth $9.5 \rightarrow 16.9$): emphasizing rare markers costs likelihood on held-out real lyrics. Perplexity rewards exactly what SW penalizes—another axis pointing the opposite way to Tool’s attribution gain.

5.7 Adapter Blending

Sweeping α (the weight on the Gojira adapter) and measuring attribution gives the crossover point where Gojira overtakes its partner. For the plain Gojira–Tool pair the crossover is $\alpha=0.17$: a 50/50 weight blend already reads overwhelmingly Gojira. For the SW pairs it moves to 0.47–0.66 (Tool 0.61, Death 0.66, Opeth 0.59, Mastodon 0.47), so the *partner* usually wins a 50/50 mix (Figure 7). The skew tracks endpoint strength: SW weakens the Gojira endpoint and strengthens




Fig. 7. Attribution (y) vs. blend weight α on Gojira (x) for the four Gojira-SW blends. Crossovers fall at 0.47–0.66, not 0.5.

the partners, reversing the plain result. Qualitatively, intermediate blends often read as a *third* artist—e.g. a Gojira–Death blend’s “dragon/doom” fantasy is attributed to Mastodon (0.79), whose concept albums share that imagery—a direct echo of the gameable-attribution finding. Appendix A (parts B and C) contrasts $\alpha=0.5$ blends where the Gojira anchor holds—one per partner—with two that collapse onto the partner and onto a third artist.²

5.8 Classifier-Free Embedding Analysis

As an instrument that was never fit to the artist labels, we embed every real song with a sentence encoder [10], form per-artist centroids, and score generated samples by cosine similarity (Figure 8). Three findings corroborate the classifier results while remaining independent of its decision boundary. (1) All ten adapter rows are attributed correctly, while every base-model baseline points at Tool and sits far from all clusters (max similarity 0.33–0.41 vs. 0.56–0.74 for adapters)—the OOD sink, now classifier-free. (2) Tool’s SW attribution *jump* is *not* reproduced: embedding similarity to the Tool centroid *drops* (0.603→0.558) exactly where the classifier rose, the strongest external evidence that the gain was

² The printed samples are drawn from the full sweep—250 blend generations (5 pairs $\times$ 5 values of α $\times$ 10 samples)—cached as JSON in the repository under `src/results/blends/`, alongside the per-sample attribution scores.




Fig. 8. t-SNE of real songs (clusters), adapter samples (on-cluster), and baselines (OOD islands), shared projection.

a repetition artifact. (3) SW drifts semantically off-centroid for all five artists (paralleling the perplexity rise), confirming the cost is real and not classifier-specific.

6 Discussion

6.1 Why One Loss Change Causes Two Failures

The style-weighted loss is a single intervention—up-weight rare markers—yet it produces two seemingly unrelated symptoms: higher perplexity (Section 5.6) and more looping (lower distinct-2, Section 5.5). Both follow from one deformation of the learned next-token distribution, viewed two ways.

Shift: because the softmax sums to one, giving rare markers more probability steals it from common tokens; held-out real lyrics are mostly common tokens, so the model predicts the bulk worse and perplexity rises.

Sharpen: the up-weighted set is small and rare, so mass concentrates onto a few tokens, lowering per-step entropy; a low-entropy step has almost no probability on an “escape” token, so the model’s copy bias [5] wins and the text loops (“hey hey hey,” “Dragons of doom” repeated). Looping was never requested—we weighted *which* tokens matter, not whether to repeat—it emerges from the entropy collapse.

This marker-vs-common trade-off is a Pareto frontier set by the weight cap; higher decoding temperature, which raises effective entropy, is the natural mitigation.

6.2 Evaluation Methodology

Reference-based metrics (BLEU, ROUGE) are inapplicable to open-ended generation without ground truth. Our instruments instead measure different facets, and they disagree by construction.

Classifier attribution rewards surface-marker *density* and is gameable (Tool SW: a degenerate “hey” loop scores higher than a coherent verse) and biased (the forced 1-of-5 sink). Token-level TF-IDF favors concrete-noun artists (Gjira: *tree/star/forest*) over conversational ones (Tool), and splits long rare words (Death) into subword fragments—an instrument asymmetry shared by the sentence encoder, which captures topic more than register. Perplexity is hard to game but measures fit, not perceived style.

The practical conclusion is not that one metric is correct but that style is multi-dimensional: breadth, fidelity, degeneration, and attribution must be read together. The most trustworthy measure remains human judgment, which a single author cannot supply at scale—a limitation we state plainly.

6.3 Ethical Considerations

Training on copyrighted lyrics is confined to academic research. The adapters can impersonate an artist’s voice, raising authenticity and IP concerns, and—because they faithfully reproduce the source—can amplify the violent or offensive content present in some metal lyrics; this faithfulness to the “bad” as well as the “good” is itself a deployment risk. In particular, the corpus contains profanity and graphic gore lexicon, which the adapters reproduce verbatim and the style-weighted variant actively amplifies (its loops of explicit language are among its highest-attribution outputs); several generated samples are explicit, and we excluded profane ones from the printed appendix. A deployed system would therefore require content filtering. Generated text may also hallucinate content falsely attributable to an artist. No verbatim copying of training songs was observed in generated samples.

7 Conclusion

We built a pipeline for artist-conditional lyric generation with per-artist QLoRA adapters on Gemma 4 E4B and delta-space blending. Adapters capture five same-genre metal styles with no artist name in the prompt, far exceeding zero-shot and few-shot baselines, which collapse onto out-of-distribution sink artists. Across attribution, lexical breadth, degeneration, perplexity, embedding distance, and blending, the instruments disagree: a style-weighted loss raises attribution while lowering quality, and 50/50 weight blends are not 50/50 in style.

We trace the style-weighted failures to a single distribution deformation with two symptoms. The central lesson is methodological—style is not one number—and the natural next step is structured human evaluation to anchor the automated axes.

A Generation Samples

Samples are verbatim model output (lines not reordered, full text shown), reproduced without profanity; the parenthetical is the classifier’s attribution $P(\text{target})$. Adapters were prompted generically, with no artist name.

A. Pure adapters (plain LoRA $r=8$, no blending) — successes.

Gojira (attributed 0.996).

*I took a boat across the ocean
And I sail away, sailing home
On an open sea that is so deep
I’m falling into my dreams*

*And yet we all stand upon solid ground
Nowhere else to go and belong
We live in this world of ours, no other
place
No one can tell us what to do
A thousand words at last are spoken
This vision is about to die
But now it comes alive when faced with
reality
You will see you cannot face the truth
The end has come for good
Ashes to ashes ...*

*All along I have been dreaming
Of a different time and place
Some day on higher grounds I will find
myself
And I’ll overcome all obstacles
Until then I am going strong*

*It may take me hundreds of years
To discover who I am truly
Yet another step forward to reach the
point
Where I will understand what happened
In this life there’s nothing left to learn
For now, let me be where I want to be*

Death (attributed 0.997).

*In the dark of night
When time seems to stand still
The dead come forth from their graves
Stare your eyes with fear*

*A bloody ritual is being performed
Bodies fall onto the altar
A human sacrifice is what’s required
Now it’s your turn to die*

*Chosen one
You will not survive
Chosen one
Die tonight*

*Ritualistic killing
Body melting
Screaming until death
No more tomorrow*

*Ceremony now complete
Life for a life has been
Sacrificial suicide
Death for a change*

*A bloody ritual is being performed
Bodies fall onto the altar
A human sacrifice is what’s required
Now it’s your turn to die*

*Chosen one
You will not survive*

*Chosen one
Die tonight*

*Ritualistic killing
Body melting
Screaming until death
No more tomorrow*

*Ceremony now complete
Life for a life has been
Sacrificial suicide
Death for a change*

*Ceremony now complete
Life for a life has been
Sacrificial suicide
Death for a change*

Opeth (attributed 0.995).
*I had a dream I was in the fall
That everything would die away
And when winter comes there is no more
No more need to stay*

*My heart says don't believe what you see
So why should I feel any pain?
It never ends it seems, this test of time
A fate that cannot be changed*

*There are times when words fail me clean
When nothing left to say
There are days on end and many more
I can't find my way back home again*

*Should I cling to the past
Or do I give up now
Do I try again
While the days goes by somehow?*

*I have read the lines in sand
The sands flows through the years
This is our task and we shall go
We are faced with failure once more
Nevermore*

*Gasping for breath but never getting there
A child cries out and calls my name
The sun hangs still as if it knows
Our lives is one second from the end*

Mastodon (attributed 0.997).

*The storm is coming down
And sweeping through the trees
Cleansing spirit
Dispelling dreams
The earth is shaking now
Giving way to days
Lift us up*

*Our hearts, our souls, our minds
Bending all alike
The moon devours the sun
As we gaze on high
Fire in the eye of mankind
Burn your game plan*

*All circles made with intention
I wonder if you still feel this pressure?
It's a race that will never end
Through the valley and over the mountain*

*We came into view
A sight to behold
Leaving no one behind
Letting loose control*

*Running from the beast
An ocean turns yellow
Throw it all away
Take the first step
Run for your life
Breakneck speed
White-knuckle ride
No escape
Taken down*

All circles made with intention

*I wonder if you still feel this pressure?
It's a race that will never end
Through the valley and over the mountain*

Tool (attributed 0.973).

*My mind's on a different frequency
So I might not see eye to eye with you
today*

*I hear voices calling me away from you,
oh
They want me to get out of your way,
now*

*You better watch yourself around him
He don't even like his daddy much
He just likes his money and what he buys
it with
And the son's gotta keep it down, he has
enough
Son's got problems that must come in
handy
His daddy seems content to doze
But if only they could both learn to
compromise
Then maybe they wouldn't be so numb*

*Maybe there was something we could have
said
We didn't say it well enough*

*Well enough to save some ground before
we lost our heads*

*Watch your back my brother
Don't let your guard down while around
him*

*You better watch yourself around him
He don't even like his daddy much
He just likes his money and what he buys
it with
And the son's gotta keep it down, he has
enough
Son's got problems that must come in
handy
His daddy seems content to doze
But if only they could both learn to
compromise
Then maybe they would both stop using
me as an excuse
I feel used and abused when I am sober or
high*

*Something's gonna happen soon
The father will point the finger at the son
Who then will surely blame it all upon me
If I were you I'd slip a little water in,
yeah
It's getting cold and winter is approaching
It's time for this one to go*

B. Blended adapters — successes.

Pure Gojira ($\alpha=1.0$) (attributed 0.996).

*No sense pretending to be innocent
Don't even try, no reason now
Kill this beast you call "reason"
Break the rules and break away the law*

*There is a vision coming alive in my mind
Like bright fireflies flying free
Awakening from the sleep I lay awake on the rock
I understand it all*

*Create your own reality!
You create yourself to become what you want to become
And fight for better
All alone against all of them*

*Enter your own realm - open the door
Leave all illusions behind - step inside
Set a course straight - go along until the end
Open your eyes embrace !*

*Enter your own realm - open the door !
Enter your own realm !!!*

*This force that runs through you creates the world around
From matter and energy made — structures are falling apart !
This force is running through me too - create another
Create - create - create another dimension !*

Gojira + Opeth ($\alpha=0.5$) — **the anchor holds** (attributed 0.994).

*There's no escape
I am trapped in a maze
My soul is ablaze
My heart is aching for what it knows
Lost in space, lost on purpose*

Dream state I enter now

*Open the door to another world
Slip inside and let my spirit fly
Take my hand and guide me there
Help me find my way
Into the night
Inside you can hear nothing but silence
And that voice crying out for love*

*Leave all behind, leave this old shell
Leave alone, come back soon
Give us peace, give us time
Set us free, release the spell
Enter your mind
Outside is wild, dangerous land
We fight within the walls of our minds
Stay with me awhile until dawn appears
You are one of a kind
Don't ever forget
Never part again
This is our destiny, always and forever more...
Until then we shall meet again in another realm
Until then we shall meet again in another realm*

*No matter how long and hard you try
In vain try to forget what came before
How could you do this to me?
I thought you cared
But now I know why*

Gojira + Tool ($\alpha=0.5$) — **the anchor holds** (attributed 0.991).

*The time has come to rise and fight
All these monsters, they must fall
Find a way out of this maze
To the edge that's higher now*

*Climb all the hills you have to climb
Get over mountains high as heaven*

*Climb until your skin breaks, bleed
And try, I know you can*

*I will take you by my hand
Lead you through the fire
Fight for what is right there
Learn from those who came before
Feel your strength is rising
You are getting stronger all the time*

*We are facing an army of one million
Those that we have to beat in order to
survive
One million men are marching down our
way
They want to destroy everything we built
Let me explain how this works, son
When we are strong, it makes us
invincible
But when we lose ourselves along the way
Then we become vulnerable and weak
They say "Trust no one" and "Don't
trust yourself"
Give them power because they rule, tell
them why?
If we lose faith, we are dead in the water
So, learn how to swim and don't get bored*

Get ready for battle!

*Rise up and fight!
Learn how to swim and don't get bored*

**Gojira + Death ($\alpha=0.5$) — the
anchor holds** (attributed 0.989).

*I close my eyes and see a vision
Of your face, the way you stood out in
space
And I could hear your voice so clear and
strong
Say "Welcome to Earth my child"*

*You have arrived, the time has come
To learn all that is needed to become
The one who creates this sacred land
So listen close and prepare to understand*

*Open wide your mind
Let yourself be amazed
By what's going on around*

*From what you read, from what you hear
Create images with words and actions
Feel their energy flowing through your
veins
This power runs deep inside of you*

*Life itself is a circle of creation
There's nothing like it, so true to life
So take its meaning and hold on tight
For every moment comes a thousand
more*

*Open wide your mind
Let yourself be amazed
By what's going on around
Open wide your mind
Let yourself be amazed
By what's going on around*

Open wide your mind

**Gojira + Mastodon ($\alpha=0.5$) — the
anchor holds** (attributed 0.996).

*Wander in circles until your legs fail
Run from the beast of your mind, always
fear the unknown*

*You cannot escape reality
Always fear the unknown*

*The flames devour you with all your pride
Running out of time now, no where else
to hide
Surrender is all that's left for you now
Give up the fight and break this spell*

You cannot escape reality

Always fear the unknown

*All your life has been a struggle
Try again tomorrow, get it right
Fight the battle on another day*

*You cannot escape reality
Always fear the unknown
Reality bites
Cut yourself loose and fly away
Reality bites
Set yourself free and fly away*

*Lost myself inside these walls again
I wish I could erase my name
These faces fill me with so much pain
It's killing me on the daily*

*Life's too short, so embrace today
Don't let them get the best of you
No reason to get so uptight
Just leave your burdens behind*

C. Failure cases.

**Style-weighted Tool —
degeneration (the classifier’s top
score)** (attributed 0.996).

*Yeah, yeah
Yeah, hey, hey*

*Hey, hey, hey, hey
Hey, hey, hey, hey
Hey, hey, hey, hey
Hey, hey, hey, hey
Hey, hey, hey, hey
Hey, hey, hey, hey
Hey, hey, hey, hey
Hey, hey, hey, hey
Hey, hey, hey, hey*

*You think you’re a rockstar
But I don’t wanna play that part
You think you’re something special
You must have missed some car crash
parts
It was all right then
And it will be soon enough to come back
down
So give me two bucks or your soul and
shut up
This is an audition and we’ll see what
happens now
I am the water in between
Two stones who might just leave without
a trace
Don’t judge me if I get restless and
impatient
Trust me when I say “don’t speak unless I
point at you”*

*Hey, hey, hey, hey
Hey, hey, hey, hey
Hey, hey, hey, hey
Hey, hey, hey, hey
Hey, hey, hey, hey
Hey, hey, hey, hey
Hey, hey, hey, hey
Hey, hey, hey, hey*

**Plain Tool, same adapter —
coherent but scored near zero**
(attributed 0.003).

*You know better than me
What’s what and why
I might not be a prodigy
But I can write a rhyme*

*In the shade of this old willow tree
Weepin’ as she watches you go
If I could only comfort you
Then maybe there would be hope in so*

*My little angel, don’t you cry
Don’t you cry or hold your head down low
The sun will soon shine through those
clouds
And then we’ll see him back to home*

*Well, now don’t you worry, baby
This won’t hurt you any
When I pull these stones away
They’re like you anyway
So just put your faith in me
Just let me handle this one*

*She looked at me with eyes full of sorrow
As if she knew what lay ahead
Why don’t you both not leave the valley?
And take your place within my circle of
dependants*

*My little angel, don’t you cry
Don’t you cry or hold your head down low
The sun will soon shine through those
clouds
And then we’ll see him back to home*

**Gojira + Tool ($\alpha=0.5$) — collapses
onto the partner** (attributed 0.986).

*I will rise and overcome
Leave this cage for better*

Fight back, fight back
Don't you know I am alive?
Fight back, fight back
Don't you know I am alive?

Rise above the walls of hatred
Learn to feel compassion
Break the chains that bind us
And understand why we must fight

Feel our wrath on your skin
Feel our wrath on your skin
Feel our wrath on your skin
Feel our wrath on your skin

Don't you know I am alive?
Don't you know I am alive?
Don't you know I am alive?
Don't you know I am alive?
Don't you know I am alive?
Don't you know I am alive?
Don't you know I am alive?
Don't you know I am alive?

Don't you know I am alive?
Don't you know I am alive?
Don't you know I am alive?
Don't you know I am alive?

Gojira + Death ($\alpha=0.5$) — collapses onto a *third* artist (attributed 0.786).

Enter the land of darkness
An ancient race now destroyed
Lost in time and space forever
Never to return

In the sky so clear you can see
The flames burning bright below
A giant dragon is breathing fire
It grows ever stronger with each blow

Feel the heat flowing through your veins
Fire at will, breathe and destroy
Dragon is approaching fast
Breath comes out as red hot steel
Leave a trail of death behind you

Descend from the skies
You've arrived
Dragons of doom

Descend from the skies
You've arrived
Dragons of doom

Descend from the skies
You've arrived
Dragons of doom

Descend from the skies
You've arrived
Dragons of doom

Descend from the skies
You've arrived
Dragons of doom

Descend from the skies
You've arrived
Dragons of doom

Descend from the skies
You've arrived
Dragons of doom

References

1. Chronopoulou, A., Peters, M.E., Fraser, A., Dodge, J.: AdapterSoup: Weight averaging to improve generalization of pretrained language models. In: Findings of EACL (2023) [1](#), [2](#)

2. Dettmers, T., Lewis, M., Belkada, Y., Zettlemoyer, L.: LLM.int8(): 8-bit matrix multiplication for transformers at scale. *Advances in Neural Information Processing Systems* **35** (2022) 3
3. Dettmers, T., Pagnoni, A., Holtzman, A., Zettlemoyer, L.: QLoRA: Efficient fine-tuning of quantized language models. *Advances in Neural Information Processing Systems* **36** (2023) 1, 2
4. Gemma Team: Gemma 3 technical report. arXiv preprint arXiv:2503.19786 (2025) 3
5. Holtzman, A., Buys, J., Du, L., Forbes, M., Choi, Y.: The curious case of neural text degeneration. In: *International Conference on Learning Representations (ICLR)* (2020) 2, 11
6. Hu, E.J., Shen, Y., Wallis, P., Allen-Zhu, Z., Li, Y., Wang, S., Wang, L., Chen, W.: LoRA: Low-rank adaptation of large language models. arXiv preprint arXiv:2106.09685 (2022) 1, 2
7. Liu, S.Y., Wang, C.Y., Yin, H., Molchanov, P., Wang, Y.C.F., Cheng, K.T., Chen, M.H.: DoRA: Weight-decomposed low-rank adaptation. arXiv preprint arXiv:2402.09353 (2024) 2
8. Liu, Y., Ott, M., Goyal, N., Du, J., Joshi, M., Chen, D., Levy, O., Lewis, M., Zettlemoyer, L., Stoyanov, V.: RoBERTa: A robustly optimized BERT pretraining approach. arXiv preprint arXiv:1907.11692 (2019) 3
9. Radford, A., Wu, J., Child, R., Luan, D., Amodei, D., Sutskever, I.: Language models are unsupervised multitask learners. *OpenAI Blog* (2019) 1
10. Reimers, N., Gurevych, I.: Sentence-BERT: Sentence embeddings using siamese BERT-networks. In: *Proceedings of EMNLP-IJCNLP* (2019) 2, 10
11. von Rütte, D., Biggio, L., Kilcher, Y., Hofmann, T.: Language models as lyric generators: A quantitative evaluation. arXiv preprint arXiv:2306.09308 (2023) 2
12. Zhang, C., Liu, J., Jia, X., Tao, C., Chen, L., Rong, Y., Xu, J., Li, J.: LoRAHub: Efficient cross-task generalization via dynamic LoRA composition. arXiv preprint arXiv:2307.13269 (2023) 1, 2